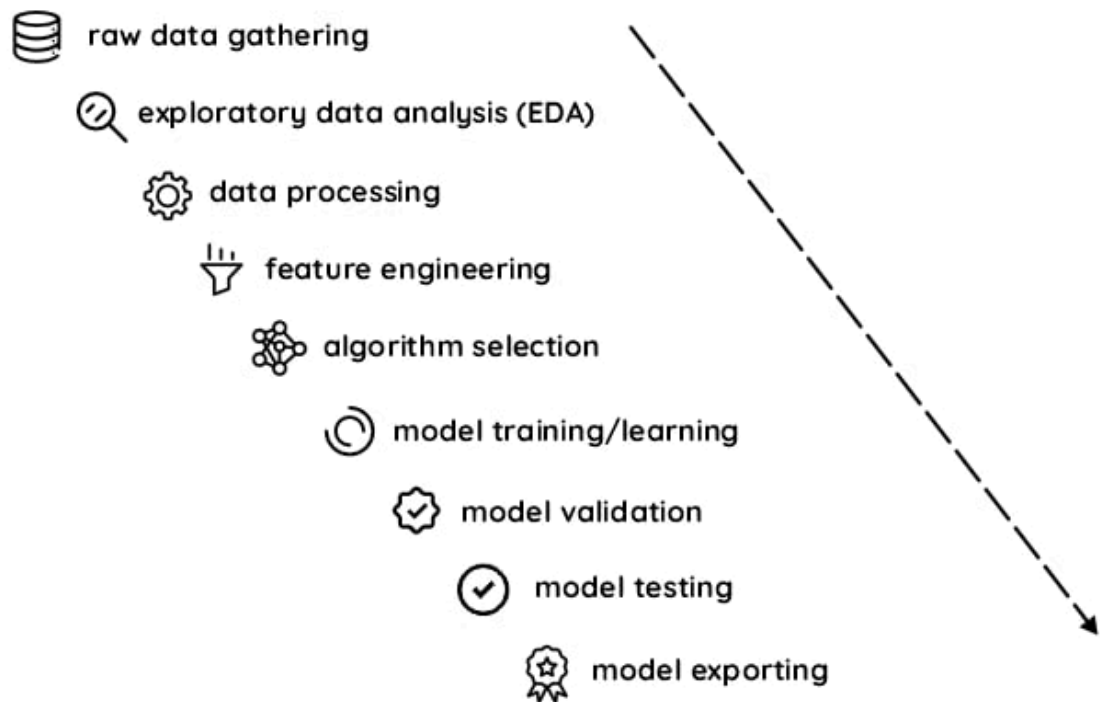


Goal: Create machine learning model that will predict quality of the wheat

Roadmap:



Loading libraries:

```
In [ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import LinearSVC
from sklearn.metrics import classification_report, accuracy_score
import joblib
from sklearn.preprocessing import LabelEncoder
```

1. Raw data gathering

```
In [2]: #dataset from: https://www.kaggle.com/datasets/muratkokludataset/durum-wh
path = "/home/darko/Desktop/durum_wheat/durum_wheat_ml/Durum_Wheat_Dataset

try:
    df = pd.read_excel(path, sheet_name=0)
    print(df.head())
except Exception as e:
    print(f"An unexpected error occurred: {e}")
```

	Target	AREA	MAJORAXIS	MINORAXIS	ECCENTRICITY	EQDIASQ	\
0	Vitreous	168	109.576157	39.396721	0.973060	213.904236	
1	Vitreous	162	105.584457	43.366894	0.964173	206.264801	
2	Vitreous	145	97.360207	35.532028	0.903072	184.619720	
3	Vitreous	178	104.080582	49.040062	1.272657	226.636627	
4	Vitreous	187	103.716667	40.885876	-1.469139	238.095779	

	PERIMETER	SOLIDITY	ROUNDNESS	SHAPEFACTOR	...	Gabor_Y9(XYZ)	\
0	202.794052	0.087454	0.017815	0.051334	...	1.187631	
1	194.794052	0.081901	0.018502	0.053650	...	2.437782	
2	177.722961	0.096026	0.019477	0.057689	...	0.527286	
3	202.509750	0.080507	0.020921	0.054543	...	2.277434	
4	198.811234	0.097650	0.022134	0.059452	...	2.337524	

	Gabor_Z1(XYZ)	Gabor_Z2(XYZ)	Gabor_Z3(XYZ)	Gabor_Z4(XYZ)	Gabor_Z5(XY
0	0.003489	0.013901	0.031153	0.003591	0.0143
1	0.003794	0.015126	0.033886	0.004166	0.0169
2	0.003069	0.012269	0.027538	0.003092	0.0124
3	0.003743	0.014940	0.033508	0.003966	0.0155
4	0.003599	0.014423	0.032525	0.003922	0.0150

	Gabor_Z6(XYZ)	Gabor_Z7(XYZ)	Gabor_Z8(XYZ)	Gabor_Z9(XYZ)
0	0.032405	0.235286	0.651593	0.775182
1	0.038586	0.368535	1.074773	1.577847
2	0.028380	0.111476	0.191517	0.350966
3	0.034776	0.180135	0.887808	1.446794
4	0.034003	0.199192	0.928664	1.786701

[5 rows x 237 columns]

2. EDA

```
In [3]: print("Dataset shape (rows, columns):", df.shape)
print('*' * 40)

print("\nDataset info:")
df.info()
print('*' * 40)

df.describe()

print('*' * 40)
```

```
print("Missing values per column:")
print(df.isna().sum())

print('*' * 40)
print("Number of duplicate rows:", df.duplicated().sum())

print('*' * 40)
sample = df.sample(random_state=42, n=5)
print("Random sample of 5 rows:\n", sample)
```

Dataset shape (rows, columns): (9000, 237)

Dataset info:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 9000 entries, 0 to 8999

Columns: 237 entries, Target to Gabor_Z9(XYZ)

dtypes: float64(235), int64(1), object(1)

memory usage: 16.3+ MB

Missing values per column:

Target 0

AREA 0

MAJORAXIS 0

MINORAXIS 0

ECCENTRICITY 0

..

Gabor_Z5(XYZ) 0

Gabor_Z6(XYZ) 0

Gabor_Z7(XYZ) 0

Gabor_Z8(XYZ) 0

Gabor_Z9(XYZ) 0

Length: 237, dtype: int64

Number of duplicate rows: 0

Random sample of 5 rows:

	Target	AREA	MAJORAXIS	MINORAXIS	ECCENTRICITY	EQDIASQ	\
7940	Foreign	113	72.095627	31.201622	0.606692	143.876068	
1162	Vitreous	154	101.589783	38.577065	0.906564	196.078888	
582	Vitreous	201	108.948685	40.072151	-0.972156	255.921143	
4081	Starchy	164	93.014153	43.781296	0.127331	208.811279	
8412	Foreign	136	86.967949	28.108158	-1.063594	173.160568	

	PERIMETER	SOLIDITY	ROUNDNESS	SHAPEFACTOR	...	Gabor_Y9(XYZ)	\
7940	134.610229	0.113682	0.027680	0.078367	...	1.923140	
1162	187.137177	0.090909	0.018999	0.055260	...	1.086189	
582	206.936188	0.101721	0.021561	0.058984	...	1.710551	
4081	180.710724	0.092655	0.024135	0.063108	...	0.852951	
8412	163.338165	0.126394	0.022894	0.064058	...	3.202615	

	Gabor_Z1(XYZ)	Gabor_Z2(XYZ)	Gabor_Z3(XYZ)	Gabor_Z4(XYZ)	\
7940	0.010843	0.043354	0.097381	0.011022	
1162	0.003581	0.014349	0.032340	0.003637	
582	0.004272	0.017132	0.038665	0.004399	
4081	0.004487	0.017879	0.040034	0.004561	
8412	0.019331	0.077476	0.174817	0.019529	

	Gabor_Z5(XYZ)	Gabor_Z6(XYZ)	Gabor_Z7(XYZ)	Gabor_Z8(XYZ)	\
7940	0.044028	0.099094	0.488496	1.024959	
1162	0.014513	0.032752	0.228496	0.429636	
582	0.017277	0.038844	0.207177	1.108145	
4081	0.018653	0.043123	0.224375	0.412812	
8412	0.077734	0.175071	0.692397	3.212104	

	Gabor_Z9(XYZ)
7940	1.716980
1162	0.692616
582	1.099306

4081	0.481957
8412	3.325619

[5 rows x 237 columns]

Columns explained:

Target: main variable indicating wheat quality

- 'Vitreous' : high-quality, hard durum wheat
- 'Starchy' : softer wheat
- 'Foreign' : non-durum or defective grains

Morphological features (shape and size):

- AREA : surface area of the grain
- MAJORAXIS : length of the major axis
- MINORAXIS : length of the minor axis
- ECCENTRICITY: measure of how elongated the grain is (0 = circle, 1 = line)
- PERIMETER : perimeter of the grain
- ROUNDNESS : roundness measure
- SOLIDITY : compactness of the grain shape
- SHAPEFACTOR : general shape descriptor
- EQDIASQ : equivalent diameter squared (derived from AREA)

Textural features (from Gabor filters on images):

- Gabor_X1(XYZ) ... Gabor_X9(XYZ) : horizontal/vertical orientation texture descriptors
- Gabor_Y1(XYZ) ... Gabor_Y9(XYZ) : diagonal texture descriptors
- Gabor_Z1(XYZ) ... Gabor_Z9(XYZ) : high-frequency/energy texture descriptors

Notes:

- All numerical features can be used for ML models
- Target is categorical; use LabelEncoder or one-hot encoding
- Gabor features help capture subtle differences in surface texture, important for distinguishing wheat quality

Category count


```
In [4]: category_counts = df['Target'].value_counts()

print("Category distribution (counts):")
print(category_counts)
```

```
Category distribution (counts):
Target
Vitreous    3000
Starchy     3000
Foreign     3000
Name: count, dtype: int64
```

Category visualisation

```
In [5]: plt.figure(figsize=(22, 5))
plt.bar(category_counts.index, category_counts.values, color=['skyblue',
plt.title("Category Distribution")
plt.xlabel("Category")
plt.ylabel("Wheat Count")
plt.show()
```



3. Data processing (cleaning)

```
In [6]: def standardize_data(df):
df.columns = df.columns.str.lower().str.replace(' ', '_').str.replace
return df

def prepare_data(df):
return standardize_data(df)
```

Data after cleaning:

```
In [7]: df = prepare_data(df)

print(df.dtypes)
```

```
target          object
area            int64
majoraxis       float64
minoraxis       float64
eccentricity    float64
...
gabor_z5xyz     float64
gabor_z6xyz     float64
gabor_z7xyz     float64
gabor_z8xyz     float64
gabor_z9xyz     float64
Length: 237, dtype: object
```

4. Feature engineering

Deciding which columns are relevant for our future model...

```
In [8]: le = LabelEncoder()
df['target_encoded'] = le.fit_transform(df['target'])

correlations = df.corr(numeric_only=True)['target_encoded'].sort_values(ascending=False)
print(correlations)
```

```
target_encoded    1.000000
gabor_a*3         0.774117
gabor_a*2         0.773761
gabor_a*1         0.773103
mean_a*           0.772723
...
gabor_b6         -0.727447
mean_b           -0.729039
gabor_b1         -0.729089
gabor_b2         -0.729253
gabor_b3         -0.729451
Name: target_encoded, Length: 237, dtype: float64
```

Separating important columns (with correlation more than 0.3)

```
In [9]: important_features = correlations[correlations.abs() > 0.3].index.tolist()
print("Important features based on correlation threshold:\n", important_features)

important_features_without_column_target = [f for f in important_features if f != 'target_encoded']
print('Number of columns in important_features_without_column_target:', len(important_features_without_column_target))
```

Important features based on correlation threshold:

```
['target_encoded', 'gabor_a*3', 'gabor_a*2', 'gabor_a*1', 'mean_a*', 'gabor_a*6', 'gabor_a*4', 'gabor_a*5', 'db4_a*', 'gabor_cr3', 'gabor_cr2', 'gabor_cr1', 'mean_cr', 'gabor_cr6', 'gabor_cr5', 'gabor_cr4', 'db4_cr', 'skew_v', 'skew_r', 'skew_l', 'skew_g', 'skew_y', 'skew_xyz', 'mean_s', 'gabor_s1', 'gabor_s2', 'gabor_s3', 'gabor_s5', 'gabor_s6', 'gabor_s4', 'skew_xyz', 'db4_s', 'entropy_s', 'entropy_v', 'entropy_b', 'skew_zxyz', 'skew_b', 'skew_b*', 'entropy_g', 'kurtosis_yxyz', 'kurtosis_xxyz', 'stddev_s', 'entropyl', 'entropy_y', 'minoraxis', 'db4_b*', 'mean_b*', 'gabor_b*1', 'gabor_b*6', 'gabor_b*5', 'gabor_b*2', 'gabor_b*3', 'gabor_b*4', 'majoraxis', 'kurtosis_zxyz', 'perimeter', 'entropy_r', 'eqdiasq', 'area', 'gabor_x8xyz', 'gabor_cb3', 'gabor_cb2', 'gabor_cb1', 'roundness', 'mean_cb', 'entropy_yxyz', 'db4_cb', 'stddev_cb', 'gabor_y8xyz', 'gabor_x7xyz', 'gabor_cb5', 'gabor_cb4', 'gabor_cb6', 'gabor_cr7', 'gabor_y7xyz', 'gabor_b*7', 'entropy_b*', 'gabor_cb7', 'stddev_b*', 'gabor_cr8', 'shapefactor', 'gabor_x9xyz', 'gabor_y9xyz', 'gabor_b*8', 'gabor_cb8', 'compactness', 'entropy_xxyz', 'skew_a*', 'stddev_zxyz', 'stddev_a*', 'gabor_a*9', 'gabor_z8xyz', 'gabor_cr9', 'gabor_a*8', 'gabor_b*9', 'gabor_z7xyz', 'solidity', 'kurtosis_s', 'gabor_cb9', 'gabor_a*7', 'gabor_z9xyz', 'entropy_cr', 'entropy_a*', 'entropy_h', 'db4_r', 'gabor_r4', 'gabor_r2', 'gabor_r1', 'gabor_r3', 'mean_r', 'gabor_r5', 'gabor_r6', 'skew_cb', 'gabor_h7', 'gabor_h8', 'stddev_h', 'gabor_h9', 'entropy_zxyz', 'db4_xxyz', 'gabor_x4xyz', 'gabor_x5xyz', 'gabor_x2xyz', 'gabor_x3xyz', 'gabor_x1xyz', 'mean_xxyz', 'gabor_x6xyz', 'db4_yxyz', 'gabor_y4xyz', 'gabor_y3xyz', 'gabor_y2xyz', 'gabor_y5xyz', 'gabor_y1xyz', 'mean_yxyz', 'gabor_y6xyz', 'db4_h', 'db4_zxyz', 'mean_h', 'gabor_h1', 'gabor_h3', 'gabor_h2', 'gabor_z4xyz', 'gabor_z5xyz', 'gabor_z6xyz', 'db4_y', 'gabor_h6', 'mean_zxyz', 'gabor_z1xyz', 'gabor_z2xyz', 'gabor_r3xyz', 'gabor_h5', 'gabor_y4', 'gabor_h4', 'gabor_y2', 'gabor_y1', 'gabor_y3', 'mean_y', 'gabor_y5', 'gabor_y6', 'db4_l', 'db4_v', 'gabor_l4', 'gabor_l2', 'gabor_l1', 'gabor_l3', 'mean_l', 'gabor_l5', 'gabor_l6', 'gabor_v4', 'gabor_v2', 'gabor_v1', 'gabor_v3', 'mean_v', 'gabor_v5', 'gabor_v6', 'db4_g', 'gabor_g4', 'gabor_g2', 'gabor_g3', 'gabor_g1', 'mean_g', 'gabor_g5', 'gabor_g6', 'db4_b', 'gabor_b4', 'gabor_b5', 'gabor_b6', 'mean_b', 'gabor_b1', 'gabor_b2', 'gabor_b3']
```

Number of columns in important_features_without_column_target: 189

5: Algorithm selection:

```
In [10]: X = df[important_features_without_column_target]
y = df["target_encoded"]

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

Model selection

```
In [11]: numeric_features = X_train.columns.tolist()
preprocessor = ColumnTransformer([
    ('scaler', StandardScaler(), numeric_features)
])

models = {
    "Logistic Regression": LogisticRegression(max_iter=1000),
    "Naive Bayes": GaussianNB(),
    "Decision Tree": DecisionTreeClassifier(),
    "Random Forest": RandomForestClassifier(),
}
```



```
    "Support Vector Machine": LinearSVC(max_iter=10000)
}

for name, model in models.items():
    print(f"\nModel: {name}")

    pipeline = Pipeline([
        ('preprocessing', preprocessor),
        ('classifier', model)
    ])

    pipeline.fit(X_train, y_train)
    y_pred = pipeline.predict(X_test)

    print("Accuracy:", round(accuracy_score(y_test, y_pred), 2))
    print(classification_report(y_test, y_pred))
```

Model: Logistic Regression

Accuracy: 0.96

	precision	recall	f1-score	support
0	0.99	0.99	0.99	600
1	0.95	0.94	0.95	600
2	0.94	0.96	0.95	600
accuracy			0.96	1800
macro avg	0.96	0.96	0.96	1800
weighted avg	0.96	0.96	0.96	1800

Model: Naive Bayes

Accuracy: 0.87

	precision	recall	f1-score	support
0	0.97	0.93	0.95	600
1	0.88	0.73	0.80	600
2	0.79	0.95	0.86	600
accuracy			0.87	1800
macro avg	0.88	0.87	0.87	1800
weighted avg	0.88	0.87	0.87	1800

Model: Decision Tree

Accuracy: 0.98

	precision	recall	f1-score	support
0	0.99	1.00	1.00	600
1	0.98	0.95	0.97	600
2	0.96	0.98	0.97	600
accuracy			0.98	1800
macro avg	0.98	0.98	0.98	1800
weighted avg	0.98	0.98	0.98	1800

Model: Random Forest

Accuracy: 0.98

	precision	recall	f1-score	support
0	0.99	1.00	1.00	600
1	0.99	0.96	0.98	600
2	0.97	0.99	0.98	600
accuracy			0.98	1800
macro avg	0.99	0.98	0.98	1800
weighted avg	0.99	0.98	0.98	1800

Model: Support Vector Machine

Accuracy: 0.97

	precision	recall	f1-score	support
0	0.99	0.99	0.99	600
1	0.96	0.95	0.96	600
2	0.95	0.96	0.96	600
accuracy			0.97	1800

macro avg	0.97	0.97	0.97	1800
weighted avg	0.97	0.97	0.97	1800

Winner is Random Forest (Accuracy 0.99)

6. Model training

```
In [12]: pipeline = Pipeline([
            ("preprocessing", preprocessor),
            ("Random Forest", RandomForestClassifier())
        ])

pipeline.fit(X, y)

joblib.dump(pipeline, "../model/wheat_sorting.pkl")

print(" Model trained and saved as 'model/wheat_sorting.pkl'")

Model trained and saved as 'model/wheat_sorting.pkl'
```

7. Model testing:

```
In [15]: n_samples = 5

test_df = pd.DataFrame(
    np.random.rand(n_samples, X_train.shape[1]) * 100,
    columns=X_train.columns
)

predictions = pipeline.predict(test_df)

predicted_labels = le.inverse_transform(predictions)

# Prikaz rezultata
print("Test samples:\n", test_df.head())
print("\nPredicted wheat quality:\n", predicted_labels)
```

Test samples:

	gabor_a*3	gabor_a*2	gabor_a*1	mean_a*	gabor_a*6	gabor_a*4	\
0	18.807395	79.103037	7.982325	54.600536	62.992982	53.693457	
1	72.333991	13.344556	8.692794	56.281361	38.393546	8.534602	
2	0.262963	95.831198	54.737749	94.508253	54.024845	93.677031	
3	28.463300	77.571470	19.528751	84.624257	44.567241	20.248151	
4	91.269872	2.820933	98.469592	69.223971	29.314132	30.136178	

	gabor_a*5	db4_a*	gabor_cr3	gabor_cr2	...	gabor_g5	gabor_g6
\							
0	31.483925	24.057640	80.509177	10.648369	...	97.097530	3.390181
1	98.828300	1.118506	38.563131	23.163078	...	52.155998	44.393741
2	79.057352	28.311180	16.185022	66.309305	...	44.849881	43.302500
3	9.795285	58.482125	0.770230	87.915045	...	51.267018	28.327253
4	85.879433	77.698837	12.998159	62.197095	...	31.437740	65.929588

	db4_b	gabor_b4	gabor_b5	gabor_b6	mean_b	gabor_b1	\
0	61.673019	98.608413	18.029573	5.050664	55.514526	23.453444	
1	78.754682	3.845860	89.640629	5.399869	41.502671	54.467619	
2	52.968647	34.621883	40.945439	52.481519	97.711657	17.140010	
3	11.567126	79.055754	39.880538	43.120465	48.414567	24.243372	
4	38.505289	85.102734	64.549089	62.136134	77.314197	47.544149	

	gabor_b2	gabor_b3
0	77.998856	76.209147
1	60.539825	48.420923
2	85.121329	68.814534
3	72.113943	34.189959
4	78.355905	0.736214

[5 rows x 189 columns]

Predicted wheat quality:

['Foreign' 'Starchy' 'Foreign' 'Foreign' 'Foreign']

8. Model exporting: